

Yue_Mu_Final_Project

April 30, 2015

1 STA663 Final Project

1.1 Background

- K-means clustering method remains the most popular clustering method and one of the top 10 algorithms in data mining. However, it is not a good clustering method in terms of efficiency or quality. Recent work has focused on improving the initialization procedure, which leads to k-means++ and kmeans|| algorithm.
- Basic K-means clustering uses the Lloyd algorithm, which starts with k arbitrary centers and performs an EM-type local search till convergence. However, it is very sensitive to initialization and takes many iterations to converge.
- K-means++ solves this problem by first choosing a center at random, and then sampling one point each iteration based on their probability distribution. It does spread out the centers pretty well, but it needs k passes over the data, which can be a problem when dataset is large.
- K-means||, also called scalable K-means++, oversamples by sampling each point independently with a larger probability, which is intuitively equivalent to updating the distribution much less frequently. It turns out to be more efficient than K-means++.

1.2 Algorithm

```
In [1]: from __future__ import division
import os
import sys
import glob
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
import pandas as pd
import time
%matplotlib inline
%precision 4
plt.style.use('ggplot')
```

1.2.1 K-means Algorithm (Simple)

```
In [2]: def initialization(data,k):
        """Return the initial set of cluster centers."""

        #Generate k random indices between 0 and the number of rows in the dataset
        centers_index = np.random.choice(range(data.shape[0]), k, replace=False)
        return data[centers_index, :]

def kmeans1(centers,data,k):
```

```

"""Implement K-means clustering algorithm. Return converged cluster centers."""

converge = False
phi = []
iterations = 0
while (not converge) and (iterations < 1000):
    #Find out which cluster each data point is the closest to
    min_dist = np.zeros(data.shape[0])
    min_index = np.zeros(data.shape[0])
    for i,p in enumerate(data):
        result = minimum_distance(p,centers)
        min_dist[i] = result[0]
        min_index[i] = result[1]
    #Calculate phi
    phi_val = np.sum(min_dist)
    phi.append(phi_val)

    #Calculate the new centers
    new_centers = np.empty(centers.shape)
    for i in range(0, k):
        if data[min_index == i,:].shape[0] == 0:
            new_centers[i] = centers[i]
        else:
            new_centers[i] = np.mean(data[min_index == i, :], axis=0)

    #Compare old centers with new centers to see if the algorithm has converged
    if np.array_equal(centers, new_centers):
        converge = True
    else:
        centers = new_centers

    iterations += 1

return (iterations, phi, centers, min_index)

def minimum_distance(p, centers):
    """Return the distance to the closest cluster center"""

    min_index = 0
    min_dist = sys.float_info.max
    for i, cc in enumerate(centers):
        d = np.sum((p - cc) ** 2)
        if min_dist > d:
            min_dist = d
            min_index = i
    return (min_dist,min_index)

```

1.2.2 K-means Algorithm (Vectorized)

```

In [3]: def kmeans2(centers,data,k):
    """Implement K-means clustering algorithm. Return converged cluster centers."""

    converge = False

```

```

phi = []
iterations = 0
while (not converge) and (iterations < 1000):
    #Find the Euclidean distance between a center and a data point
    data2 = data[:, np.newaxis, :]
    d2 = (data2 - centers) ** 2
    #Calculate the total distance to each center for each data point.
    distance = np.sum(d2, axis=2)

    #Find out which cluster each data point belongs to.
    min_index = np.zeros(distance.shape)
    min_index[range(distance.shape[0]), np.argmin(distance, axis=1)] = 1

    #Calculate phi
    phi_val = np.sum(distance[min_index == 1])
    phi.append(phi_val)

    #Calculate the new centers
    new_centers = np.empty(centers.shape)
    for i in range(0, k):
        if data[min_index[:, i] == 1, :].shape[0] == 0:
            new_centers[i] = centers[i]
        else:
            new_centers[i] = np.mean(data[min_index[:, i] == True, :], axis=0)

    #Compare old centers with new centers to see if the algorithm has converged
    if np.array_equal(centers, new_centers):
        converge = True
    else:
        centers = new_centers

    iterations += 1

return (iterations, phi, centers, min_index)

```

1.2.3 K-means++ Algorithm (Simple)

```

In [4]: def kmeanspp1(data, k):
    """Implement the K-means++ algorithm. Return k initial cluster centers."""
    #Sample a point uniformly at random from the data
    centers = data[np.random.choice(range(data.shape[0]), 1), :]

    #Iterate k-1 times through the dataset to select the initial cluster centers
    while centers.shape[0] < k:
        min_dist = np.zeros(data.shape[0])
        for i, p in enumerate(data):
            min_dist[i] = minimum_distance(p, centers)[0]
        #Calculate cost of the data
        phi = np.sum(min_dist)
        #Calculate the probability distribution
        prob = min_dist/phi
        #Select the next centroid using the probability distribution calculated
        centers = np.vstack([centers, data[np.random.choice(range(data.shape[0]), 1,
                                                                p = prob), :]])

```

```
return centers
```

1.2.4 K-means++ Algorithm (Vectorized)

```
In [5]: def kmeanspp2(data, k):
    """Implement the K-means++ algorithm. Return k initial cluster centers."""
    #Sample a point uniformly at random from the data
    centers = data[np.random.choice(range(data.shape[0]),1), :]
    data2 = data[:, np.newaxis, :]

    #Iterate k-1 times through the dataset to select the initial cluster centers
    while centers.shape[0] < k:
        d2 = (data2 - centers) ** 2
        distance = np.sum(d2, axis=2)
        min_index = np.zeros(distance.shape)
        min_index[range(distance.shape[0]), np.argmin(distance, axis=1)] = 1
        min_dist=distance[min_index == 1]
        phi = np.sum(min_dist)
        #Calculate the probability distribution
        prob = min_dist/phi
        #Select the next centroid using the probability distribution calculated
        centers = np.vstack([centers, data[np.random.choice(range(data.shape[0]),1,
                                                                p = prob), :]])

    return centers
```

1.2.5 Scalable K-means++ Algorithm

```
In [6]: def scalablekmeanspp(data,k,l):
    """Implement the K-means++ algorithm. Return k initial cluster centers."""
    #Sample a point uniformly at random from the data
    centers = data[np.random.choice(range(data.shape[0]),1), :]
    data2 = data[:, np.newaxis, :]
    min_dist = np.zeros(data.shape[0])
    for i,p in enumerate(data):
        min_dist[i] = minimum_distance(p,centers)[0]
    phi = np.sum(min_dist)

    for i in range(int(round(np.log(phi)))):
        d2 = (data2 - centers) ** 2
        distance = np.sum(d2, axis=2)
        min_index = np.zeros(distance.shape)
        min_index[range(distance.shape[0]), np.argmin(distance, axis=1)] = 1
        min_dist=distance[min_index == 1]
        phi = np.sum(min_dist)
        for j,p in enumerate(data):
            prob = 1*min_dist[j]/phi
            u=np.random.uniform(0,1)
            if prob >= u:
                centers = np.vstack([centers, p])

    #Now we have an initial set of cluster centers that is greater than k
    #Find number of points in dataset closer to each center than any other centers
    d2 = (data2 - centers) ** 2
    distance = np.sum(d2, axis=2)
    min_index = np.zeros(distance.shape)
    min_index[range(distance.shape[0]), np.argmin(distance, axis=1)] = 1
```

```

weights = np.array([np.count_nonzero(min_index[:, j]) for j in
                      range(centers.shape[0])]).reshape(-1,1)
#Recluster the weighted points into k clusters using kmeans++
new_centers = centers[np.random.choice(range(centers.shape[0]),1), :]
centers2 = centers
index = np.where(centers2==new_centers)[0]
centers2 = np.delete(centers2,index[0],axis=0)
weights = np.delete(weights,index[0])
while new_centers.shape[0] < k:
    #Calculate the probability distribution based on weights
    prob = weights/np.sum(weights)
    #Select the next centroid using the probability distribution calculated
    c = centers2[np.random.choice(range(centers2.shape[0]),1, p = prob.ravel()), :]
    index = np.where(centers2==c)[0]
    new_centers = np.vstack([new_centers, c])
    #Remove the selected center and its corresponding weight
    centers2 = np.delete(centers2,index[0],axis=0)
    weights = np.delete(weights,index[0])

#Implement kmeans clustering to obtain the new centers
new_centers2 = kmeans2(new_centers,centers,k)[2]
return new_centers2

```

1.3 Simulate Data

```

In [7]: def generate_data(n):
        """Return mixed data with cluster labels."""
        mean0 = [0, 1]
        cov0 = [[1, 0.5], [0.5, 1]]
        data0 = np.random.multivariate_normal(mean0, cov0, n)
        data0 = np.hstack((data0, np.ones((data0.shape[0],1))))

        mean1 = [6, 7]
        cov1 = [[1, 0.5], [0.5, 1]]
        data1 = np.random.multivariate_normal(mean1, cov1, n)
        data1 = np.hstack((data1, np.ones((data1.shape[0],1)) * 2))

        mean2 = [11, 12]
        cov2 = [[1, 0.5], [0.5, 1]]
        data2 = np.random.multivariate_normal(mean2, cov2, n)
        data2 = np.hstack((data2, np.ones((data2.shape[0],1)) * 3))

        data = np.vstack((data0, data1, data2))
        np.random.shuffle(data)
        print (data.shape)
        return data

data1=generate_data(1000)[: ,0:2]

(3000, 3)

```

```

In [8]: def generate_data(n):
        """Return mixed data with cluster labels."""

```

```

mean0 = [0, 1]
cov0 = [[1, 0.5], [0.5, 1]]
data0 = np.random.multivariate_normal(mean0, cov0, n)
data0 = np.hstack((data0, np.ones((data0.shape[0],1))))

mean1 = [5, 6]
cov1 = [[1, 0.5], [0.5, 1]]
data1 = np.random.multivariate_normal(mean1, cov1, n)
data1 = np.hstack((data1, np.ones((data1.shape[0],1)) * 2))

mean2 = [10, 11]
cov2 = [[1, 0.5], [0.5, 1]]
data2 = np.random.multivariate_normal(mean2, cov2, n)
data2 = np.hstack((data2, np.ones((data2.shape[0],1)) * 3))

mean3 = [15, 16]
cov3 = [[1, 0.5], [0.5, 1]]
data3 = np.random.multivariate_normal(mean2, cov2, n)
data3 = np.hstack((data3, np.ones((data3.shape[0],1)) * 4))

mean4 = [20, 21]
cov4 = [[1, 0.5], [0.5, 1]]
data4 = np.random.multivariate_normal(mean2, cov2, n)
data4 = np.hstack((data4, np.ones((data4.shape[0],1)) * 5))

data = np.vstack((data0, data1, data2, data3, data4))
np.random.shuffle(data)
print (data.shape)
return data

data2=generate_data(5000)[:,:0:2]

(25000, 3)

```

1.4 Test


```

In [229]: k=3
result = kmeans2(initialization(data1,k),data1,k)
centers = result[2]
min_index = result[3]
#Plot clusters
cols=iter(['b','g','y'])
for i in range(k):
    plt.scatter(data1[min_index[:,i] == 1, :][:,0],
                data1[min_index[:,i] == 1, :][:,1], color=next(cols))
plt.scatter(centers[:,0], centers[:,1], color='r', marker='x')

('Initial Centers:', array([[ 8.2972,  8.0749],
 [ 6.5101,  6.9435],
 [-1.2449, -0.4135]]))
('Required ', 4, ' iterations to converge.')

Out[229]: <matplotlib.collections.PathCollection at 0x7fe759310e10>


```



```
In [230]: result = kmeans2(kmeanspp2(data1,k),data1,k)
          centers = result[2]
          min_index = result[3]
          #Plot clusters
          cols=iter(['b','g','y'])
          for i in range(k):
              plt.scatter(data1[min_index[:,i] == 1, :][:,0],
                          data1[min_index[:,i] == 1, :][:,1], color=next(cols))
          plt.scatter(centers[:,0], centers[:,1], color='r', marker='x')

('Initial Centers:', array([[ 10.4522,  14.5413],
 [  1.6204,   2.3635],
 [ 11.7767,  10.5528]]))
('Required ', 6, ' iterations to converge.')
```


```
Out[230]: <matplotlib.collections.PathCollection at 0x7fe7592a2e50>
```



```
In [232]: result = kmeans2(scalablekmeanspp(data1,k,2),data1,k)
centers = result[2]
min_index = result[3]
#Plot clusters
cols=iter(['b','g','y'])
for i in range(k):
    plt.scatter(data1[min_index[:,i] == 1, :][:,0],
                data1[min_index[:,i] == 1, :][:,1], color=next(cols))
plt.scatter(centers[:,0], centers[:,1], color='r', marker='x')

('Initial Centers:', array([[ 12.5395,  12.3895],
 [ 10.388 ,  12.0946],
 [  6.6019,   7.9372]]))
('Required ', 7, ' iterations to converge.')
('Initial Centers:', array([[ 11.1654,  12.1567],
 [  5.877 ,   6.9726],
 [  0.4497,   1.459 ]]))
('Required ', 3, ' iterations to converge.')

Out[232]: <matplotlib.collections.PathCollection at 0x7fe75923bf90>
```

From the output, we can see that all three algorithms are working as expected.

1.5 Comparison of Algorithms

```
In [53]: np.random.seed(200)
         result = kmeans2(initialization(data1,3),data1,3)
         #Print number of iterations till convergence
         print result[0]
         #Print the clustering cost after initialization
         print result[1][0]
```

```
7
46372.4191656
```

```
In [55]: np.random.seed(200)
         result = kmeans2(kmeanspp2(data1,3),data1,3)
         print result[0]
         print result[1][0]
```

```
4
13765.1574859
```

```
In [54]: np.random.seed(200)
         result = kmeans2(scalablekmeanspp(data1,3,1.5),data1,3)
         print result[0]
         print result[1][0]
```

```
3
6396.64504607
```

```
In [66]: np.random.seed(140)
         result = kmeans2(initialization(data2,5),data2,5)
         print result[0]
         print result[1][0]
```

```
36
179399.609486
```

```
In [65]: np.random.seed(140)
         result = kmeans2(kmeanspp2(data2,5),data2,5)
         print result[0]
         print result[1][0]
```

```
27
81570.8032694
```

```
In [64]: np.random.seed(140)
         result = kmeans2(scalablekmeanspp(data2,5,2.5),data2,5)
         print result[0]
         print result[1][0]
```

```
14
34903.9182741
```

Based on the results above, we can see that K-means++ reduces number of Lloyd iterations by a great amount, and scalable K-means++ does that even more than K-means++, especially when the dataset is bigger and has more clusters. The efficiency of scalable K-means++ can also be reflected in the initial clustering cost. The smaller the initial clustering cost is, the better the initial centroids are, meaning that they are closer to the actual cluster centers. The initial clustering cost of scalable K-means++ is significantly lower than the other two algorithms.

1.6 Code Efficiency

```
In [67]: ! pip install --pre line-profiler &> /dev/null
         ! pip install psutil &> /dev/null
         ! pip install memory_profiler &> /dev/null
```

```
In [68]: %load_ext line_profiler
```

```
In [92]: %lprun -f kmeans1 kmeans1(kmeanspp1(data1,3),data1,3)
```

```
In []: Timer unit: 1e-06 s
```

```
Total time: 0.417552 s
File: <ipython-input-2-eb941895ad31>
Function: kmeans1 at line 8
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
8					def kmeans1(centers,data,k):
9					"""Implement K-means clustering algorithm. """
10					
11	1	2	2.0	0.0	converge = False
12	1	2	2.0	0.0	phi = []
13	1	0	0.0	0.0	iterations = 0
14	4	11	2.8	0.0	while (not converge) and (iterations < 1000):

15					<i>#Find out which cluster each data point</i>
16	3	30	10.0	0.0	<code>min_dist = np.zeros(data.shape[0])</code>
17	3	10	3.3	0.0	<code>min_index = np.zeros(data.shape[0])</code>
18	9003	10552	1.2	2.5	<code>for i,p in enumerate(data):</code>
19	9000	381187	42.4	91.3	<code> result = minimum_distance(p,centers)</code>
20	9000	14599	1.6	3.5	<code> min_dist[i] = result[0]</code>
21	9000	9386	1.0	2.2	<code> min_index[i] = result[1]</code>
22					<i>#Calculate phi</i>
23	3	40	13.3	0.0	<code>phi_val = np.sum(min_dist)</code>
24	3	7	2.3	0.0	<code>phi.append(phi_val)</code>
25					
26					<i>#Calculate the new centers</i>
27	3	15	5.0	0.0	<code>new_centers = np.empty(centers.shape)</code>
28	12	32	2.7	0.0	<code>for i in range(0, k):</code>
29	9	514	57.1	0.1	<code> if data[min_index == i,:].shape[0] :</code>
30					<code> new_centers[i] = centers[i]</code>
31					<code> else:</code>
32	9	1051	116.8	0.3	<code> new_centers[i] = np.mean(data[m</code>
33					
34					<i>#Compare old centers with new centers t</i>
35	3	109	36.3	0.0	<code>if np.array_equal(centers, new_centers)</code>
36	1	1	1.0	0.0	<code> converge = True</code>
37					<code>else:</code>
38	2	1	0.5	0.0	<code> centers = new_centers</code>
39					
40	3	3	1.0	0.0	<code> iterations += 1</code>
41					
42	1	0	0.0	0.0	<code>return (iterations, phi, centers, min_index)</code>

We can see that majority of the time was spent on running the `minimum_distance` function because it requires running through the entire dataset. We can try to optimize this part of the algorithm using Cython.

```
In [9]: %load_ext cythonmagic
```

```
In [10]: %%cython -a
import numpy as np
cimport numpy as np

cimport cython
@cython.boundscheck(False)
@cython.wraparound(False)

def cminimum_distance(np.ndarray[np.float64_t, ndim=1] p,
                      np.ndarray[np.float64_t, ndim=2] centers):

    cdef int min_index = 0
    cdef float min_dist = 1.79769313486e+308
    cdef int i
    cdef np.ndarray[np.float64_t, ndim=1] cc
    for i, cc in enumerate(centers):
        d = np.sum((p - cc) ** 2)
        if min_dist > d:
            min_dist = d
            min_index = i
    return (min_dist,min_index)
```

Out[10]: <IPython.core.display.HTML at 0x7fa432851b90>

1.6.1 K-means Algorithm (Cython)

```
In [11]: def kmeans3(centers,data,k):
    """Implement K-means clustering algorithm. Return converged cluster centers."""

    converge = False
    phi = []
    iterations = 0
    while (not converge) and (iterations < 1000):
        #Find out which cluster each data point is the closest to
        min_dist = np.zeros(data.shape[0])
        min_index = np.zeros(data.shape[0])
        for i,p in enumerate(data):
            #Here use the cython version of the minimum_distance function
            result = cminimum_distance(p,centers)
            min_dist[i] = result[0]
            min_index[i] = result[1]
        #Calculate phi
        phi_val = np.sum(min_dist)
        phi.append(phi_val)

        #Calculate the new centers
        new_centers = np.empty(centers.shape)
        for i in range(0, k):
            if data[min_index == i,:].shape[0] == 0:
                new_centers[i] = centers[i]
            else:
                new_centers[i] = np.mean(data[min_index == i, :], axis=0)

        #Compare old centers with new centers to see if the algorithm has converged
        if np.array_equal(centers, new_centers):
            converge = True
        else:
            centers = new_centers

        iterations += 1

    return (iterations, phi, centers, min_index)
```

1.6.2 K-means++ (Cython)

```
In [12]: def kmeanspp3(data, k):
    """Implement the K-means++ algorithm. Return k initial cluster centers."""
    #Sample a point uniformly at random from the data
    centers = data[np.random.choice(range(data.shape[0]),1), :]

    #Iterate k-1 times through the dataset to select the initial cluster centers
    while centers.shape[0] < k:
        min_dist = np.zeros(data.shape[0])
        for i,p in enumerate(data):
            #Here use the cython version of the minimum_distance function
            min_dist[i] = cminimum_distance(p,centers)[0]
        #Calculate cost of the data
```

```

    phi = np.sum(min_dist)
    #Calculate the probability distribution
    prob = min_dist/phi
    #Select the next centroid using the probability distribution calculated
    centers = np.vstack([centers, data[np.random.choice(range(data.shape[0]),1,
                                                         p = prob), :]])

    return centers

In [13]: %timeit -n 5 kmeans1(kmeanspp1(data1,3),data1,3)
          %timeit -n 5 kmeans2(kmeanspp2(data1,3),data1,3)
          %timeit -n 5 kmeans3(kmeanspp3(data1,3),data1,3)

5 loops, best of 3: 469 ms per loop
5 loops, best of 3: 5.36 ms per loop
5 loops, best of 3: 531 ms per loop

```

It seems that in this case vectorization increases the efficiency the most.